

All Training materials are provided "as is" and without warranty and RStudio disclaims any and all express and implied warranties including without limitation the implied warranties of title, fitness for a particular purpose, merchantability and noninfringement.

Transforming Data

Turn data into information
with dplyr and base R



Garrett Grolmund

Master Instructor, RStudio

August 2014

1. Baby names data
2. Subsetting data sets
3. Transform and reorder data
4. Join data sets
5. Groupwise operations

Baby names data

Baby names

Top 1000 male and female baby names in the US, from 1880 to 2008.

258,000 records ($1000 * 2 * 129$)

But only five variables: year, name, soundex, sex and prop.


```
library(ggplot2)
```

Make sure the data
sets are in your
working directory

Prevent R from
reading in strings as
factors (the default)

```
options(stringsAsFactors = FALSE)
```

```
bnames <- read.csv("data/bnames.csv.bz2")
```

```
births <- read.csv("data/births.csv")
```


head(bnames)

	year	name	prop	sex	soundex
1	1880	John	0.081541	boy	J500
2	1880	William	0.080511	boy	W450
3	1880	James	0.050057	boy	J520
4	1880	Charles	0.045167	boy	C642
5	1880	George	0.043292	boy	G620

tail(bnames)

	year	name	prop	sex	soundex
257996	2008	Carleigh	0.000128	girl	C642
257997	2008	Iyana	0.000128	girl	I500
257998	2008	Kenley	0.000127	girl	K540
257999	2008	Sloane	0.000127	girl	S450
258000	2008	Elianna	0.000127	girl	E450

Popularity of Garrett, Hadley and Winston (1880–2008)



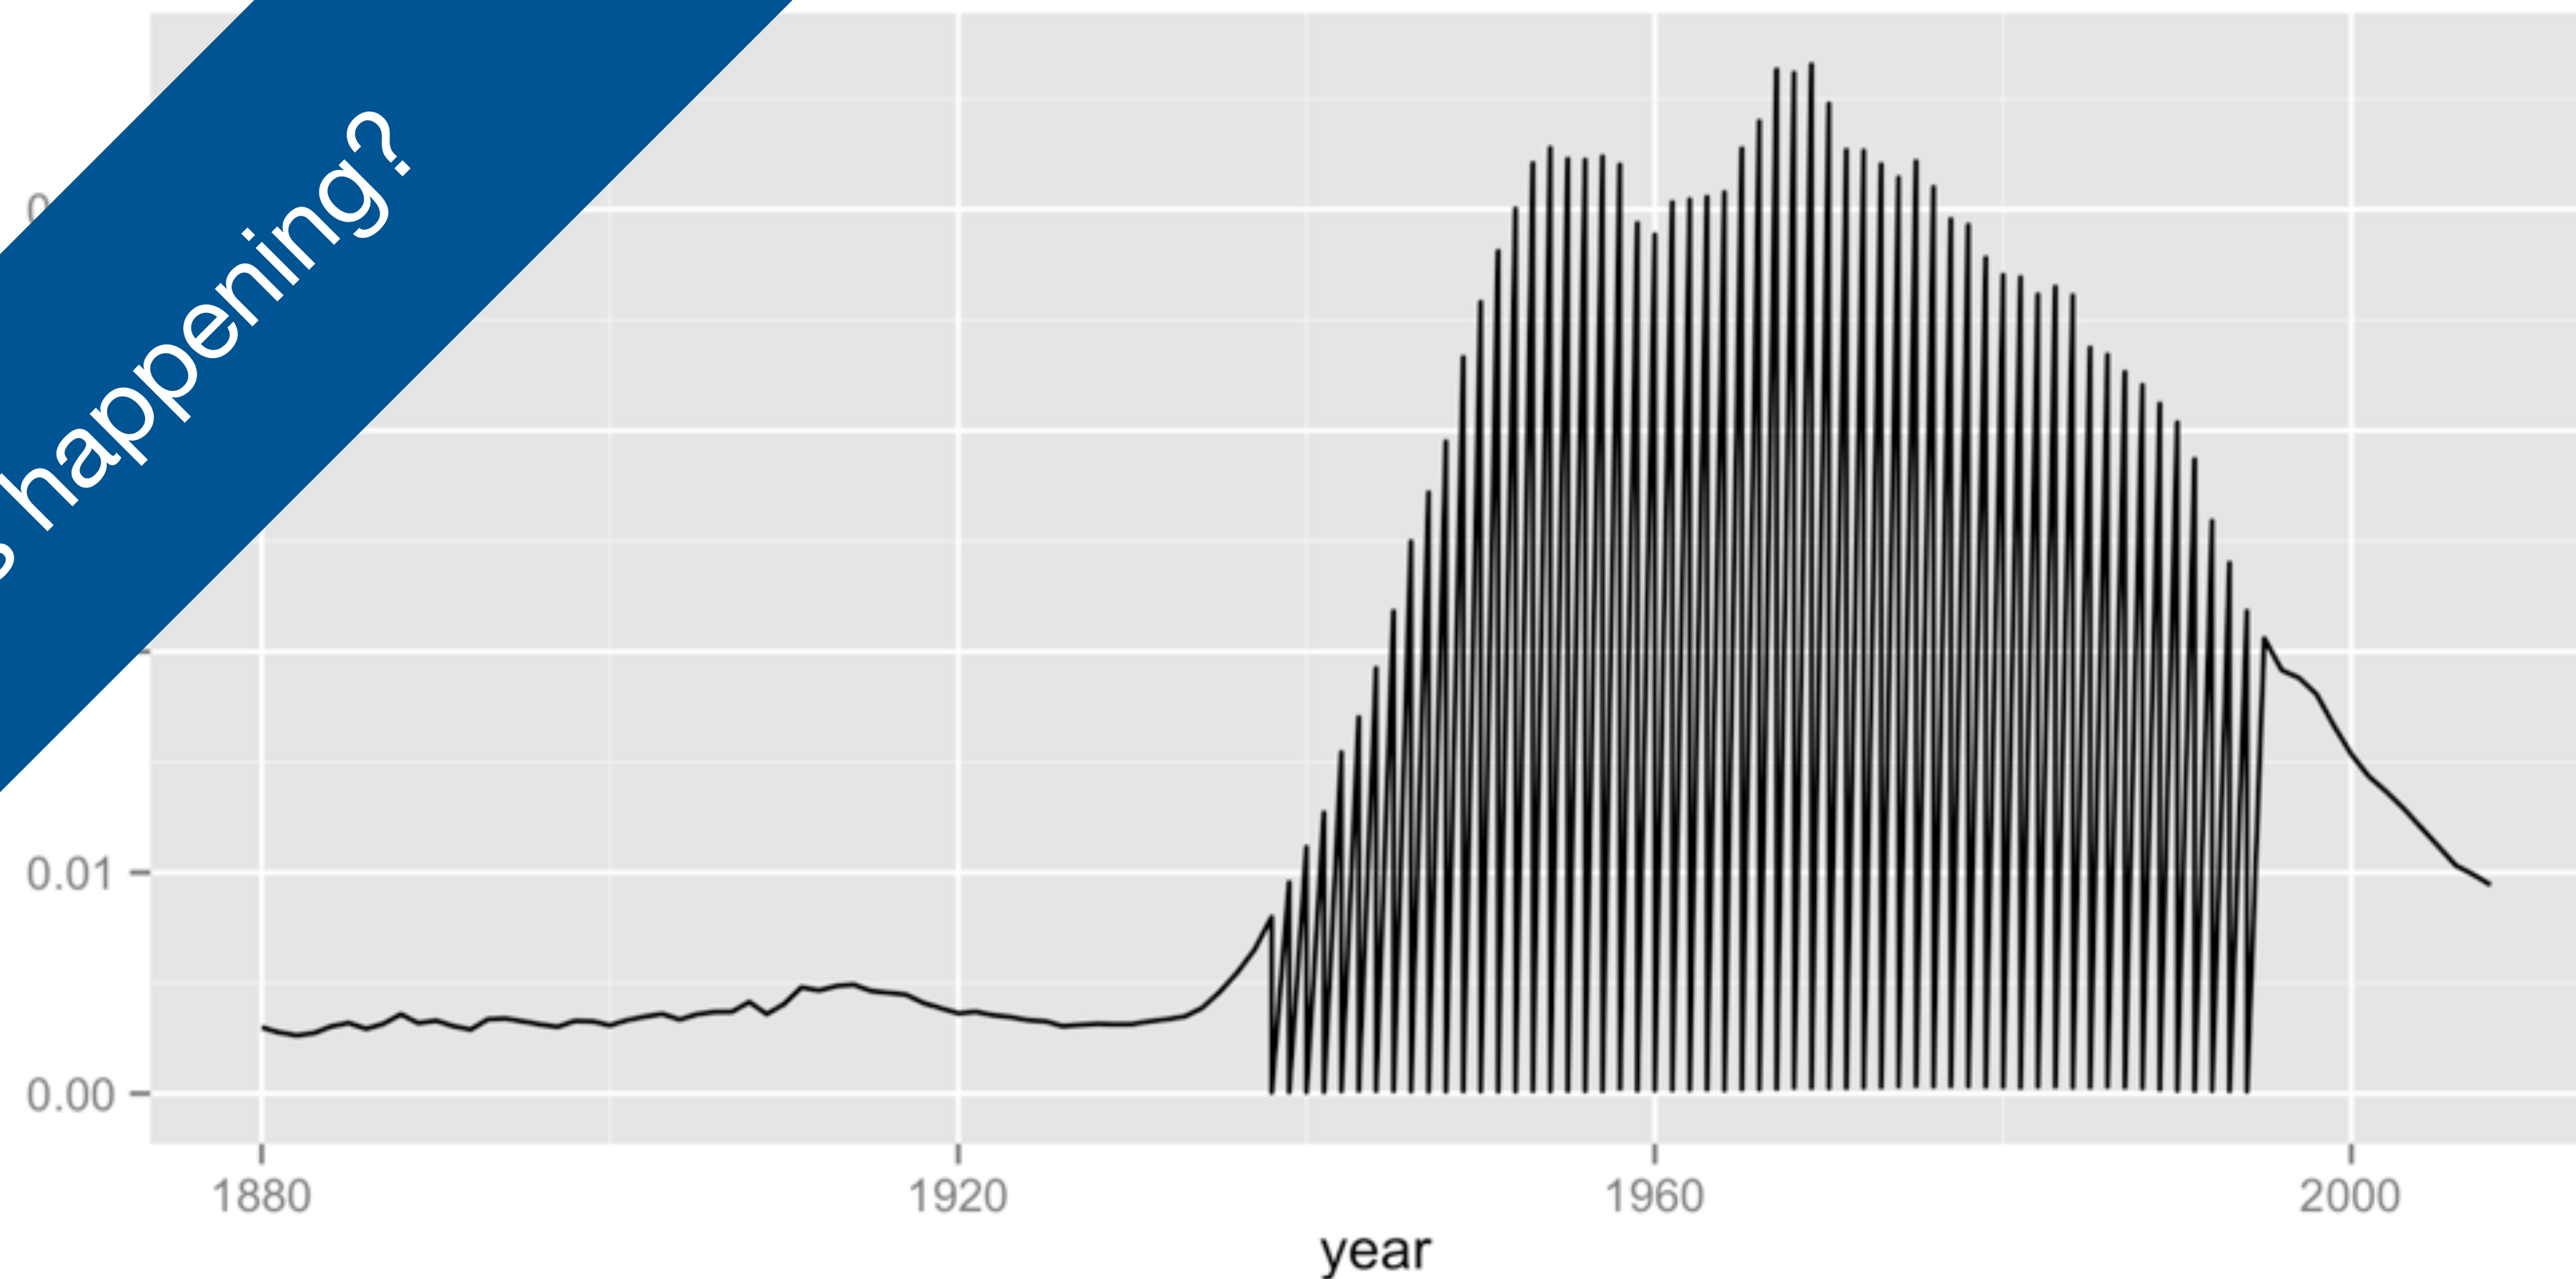
Your turn

Use logical subsetting to extract your name from the dataset. Plot the trend over time.

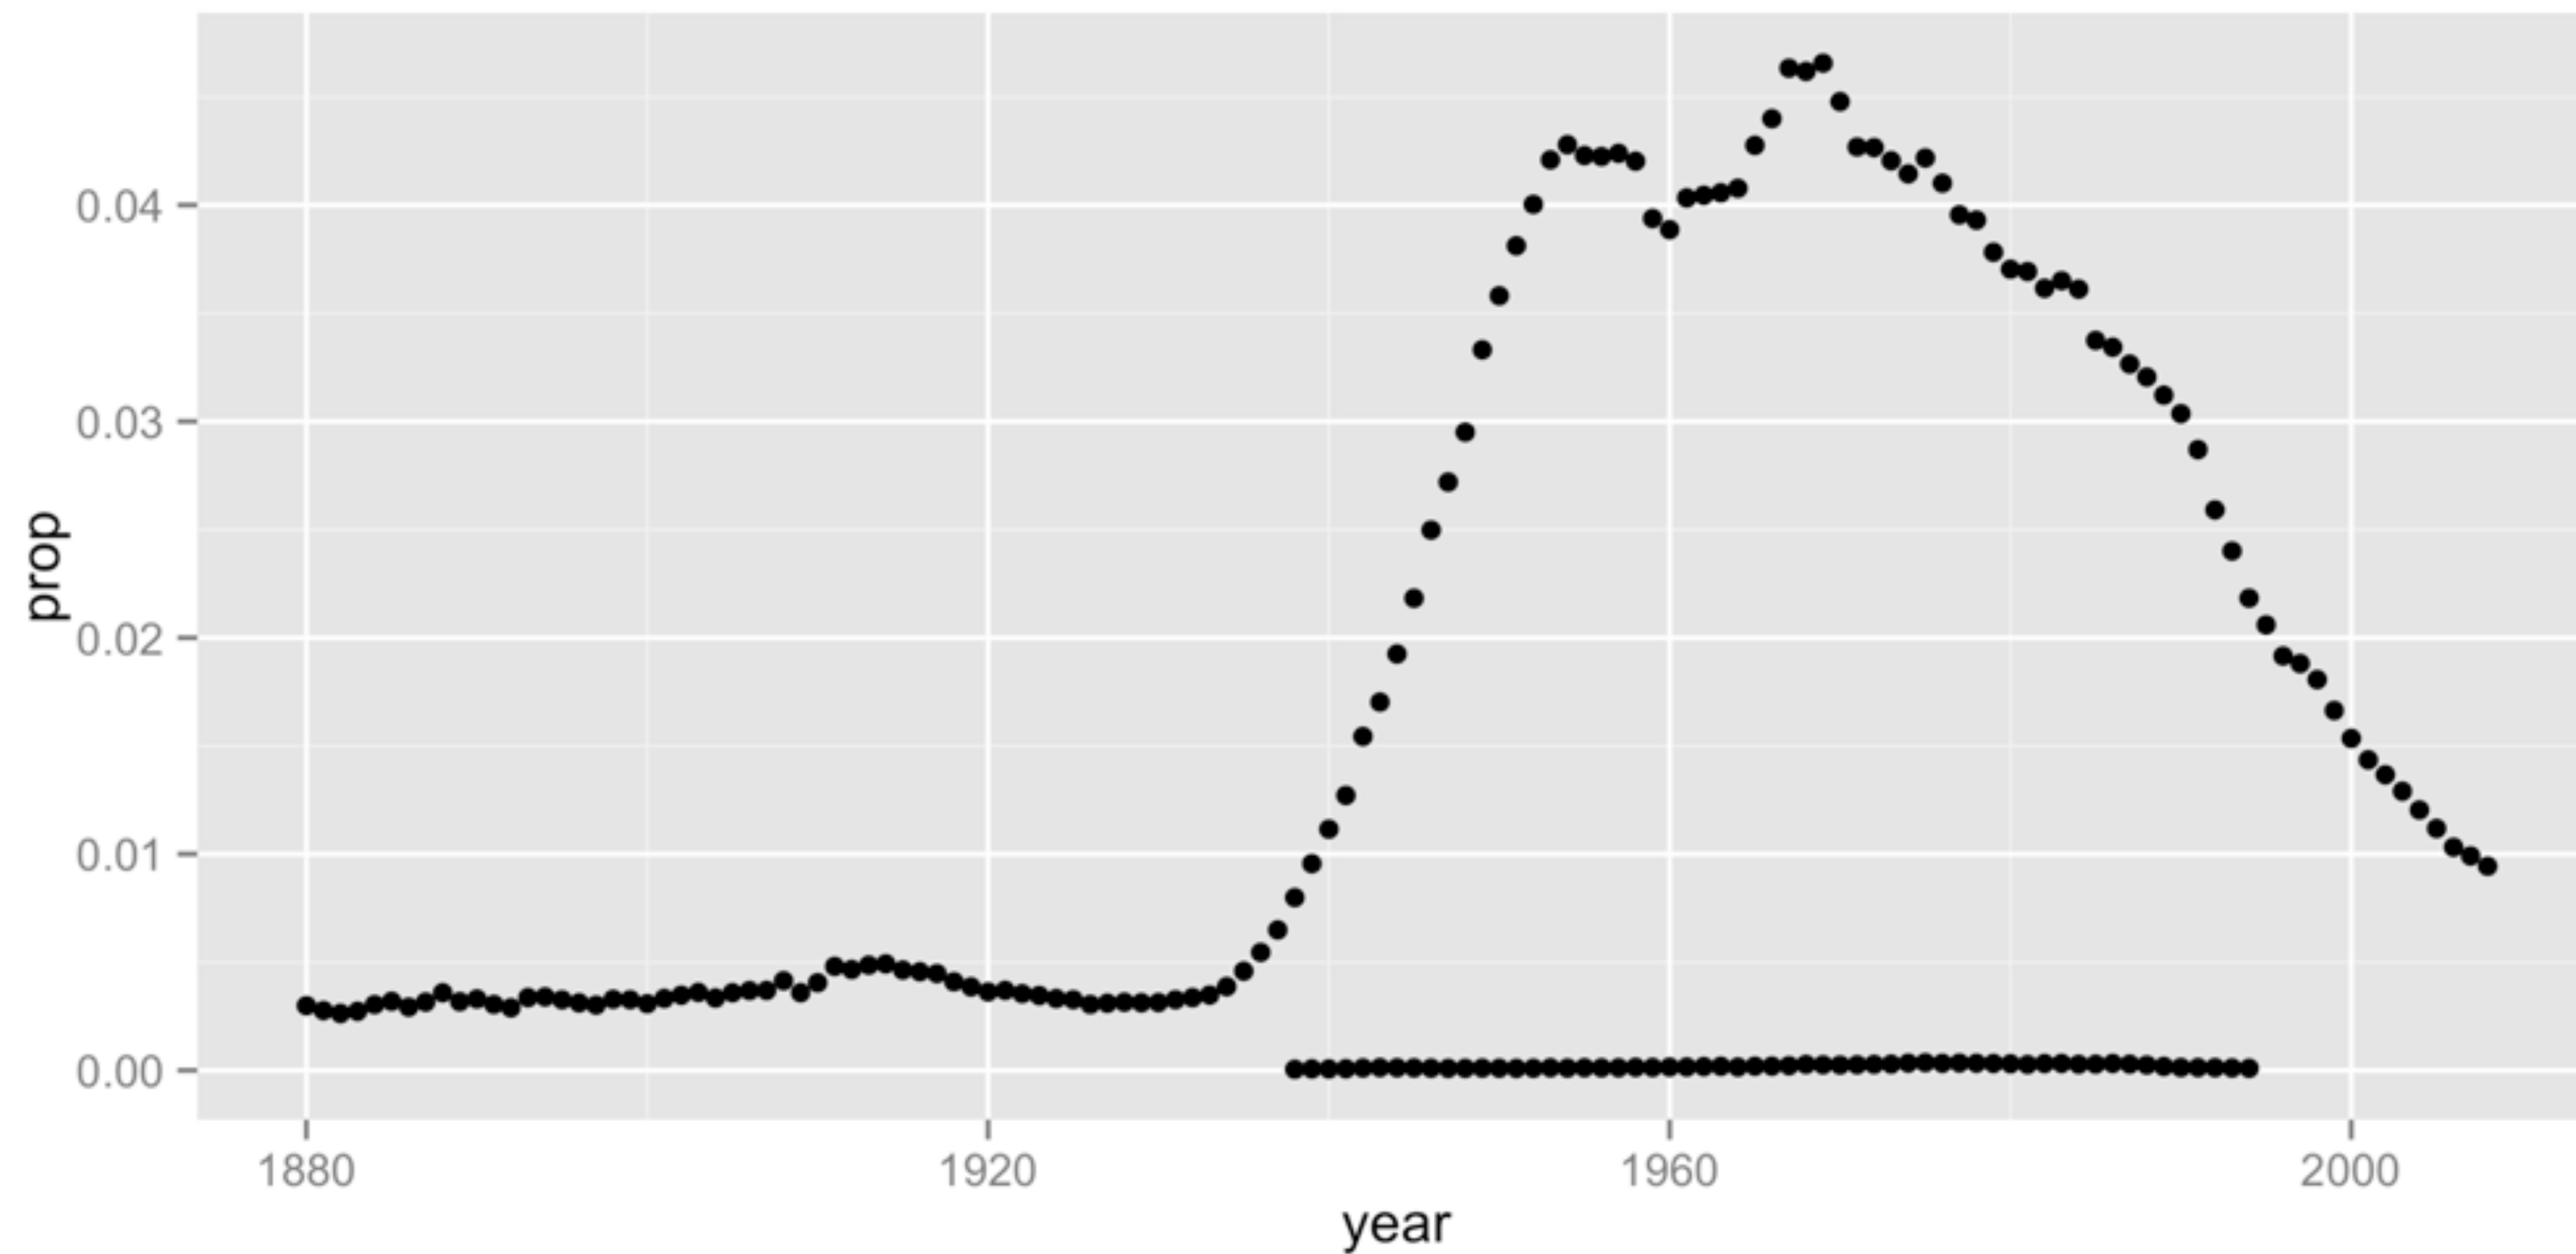
What geom should you use? Do you need any extra aesthetics?

```
garrett <- bnames[bnames$name == "Garrett", ]  
qplot(year, prop, data = garrett, geom = "line")
```

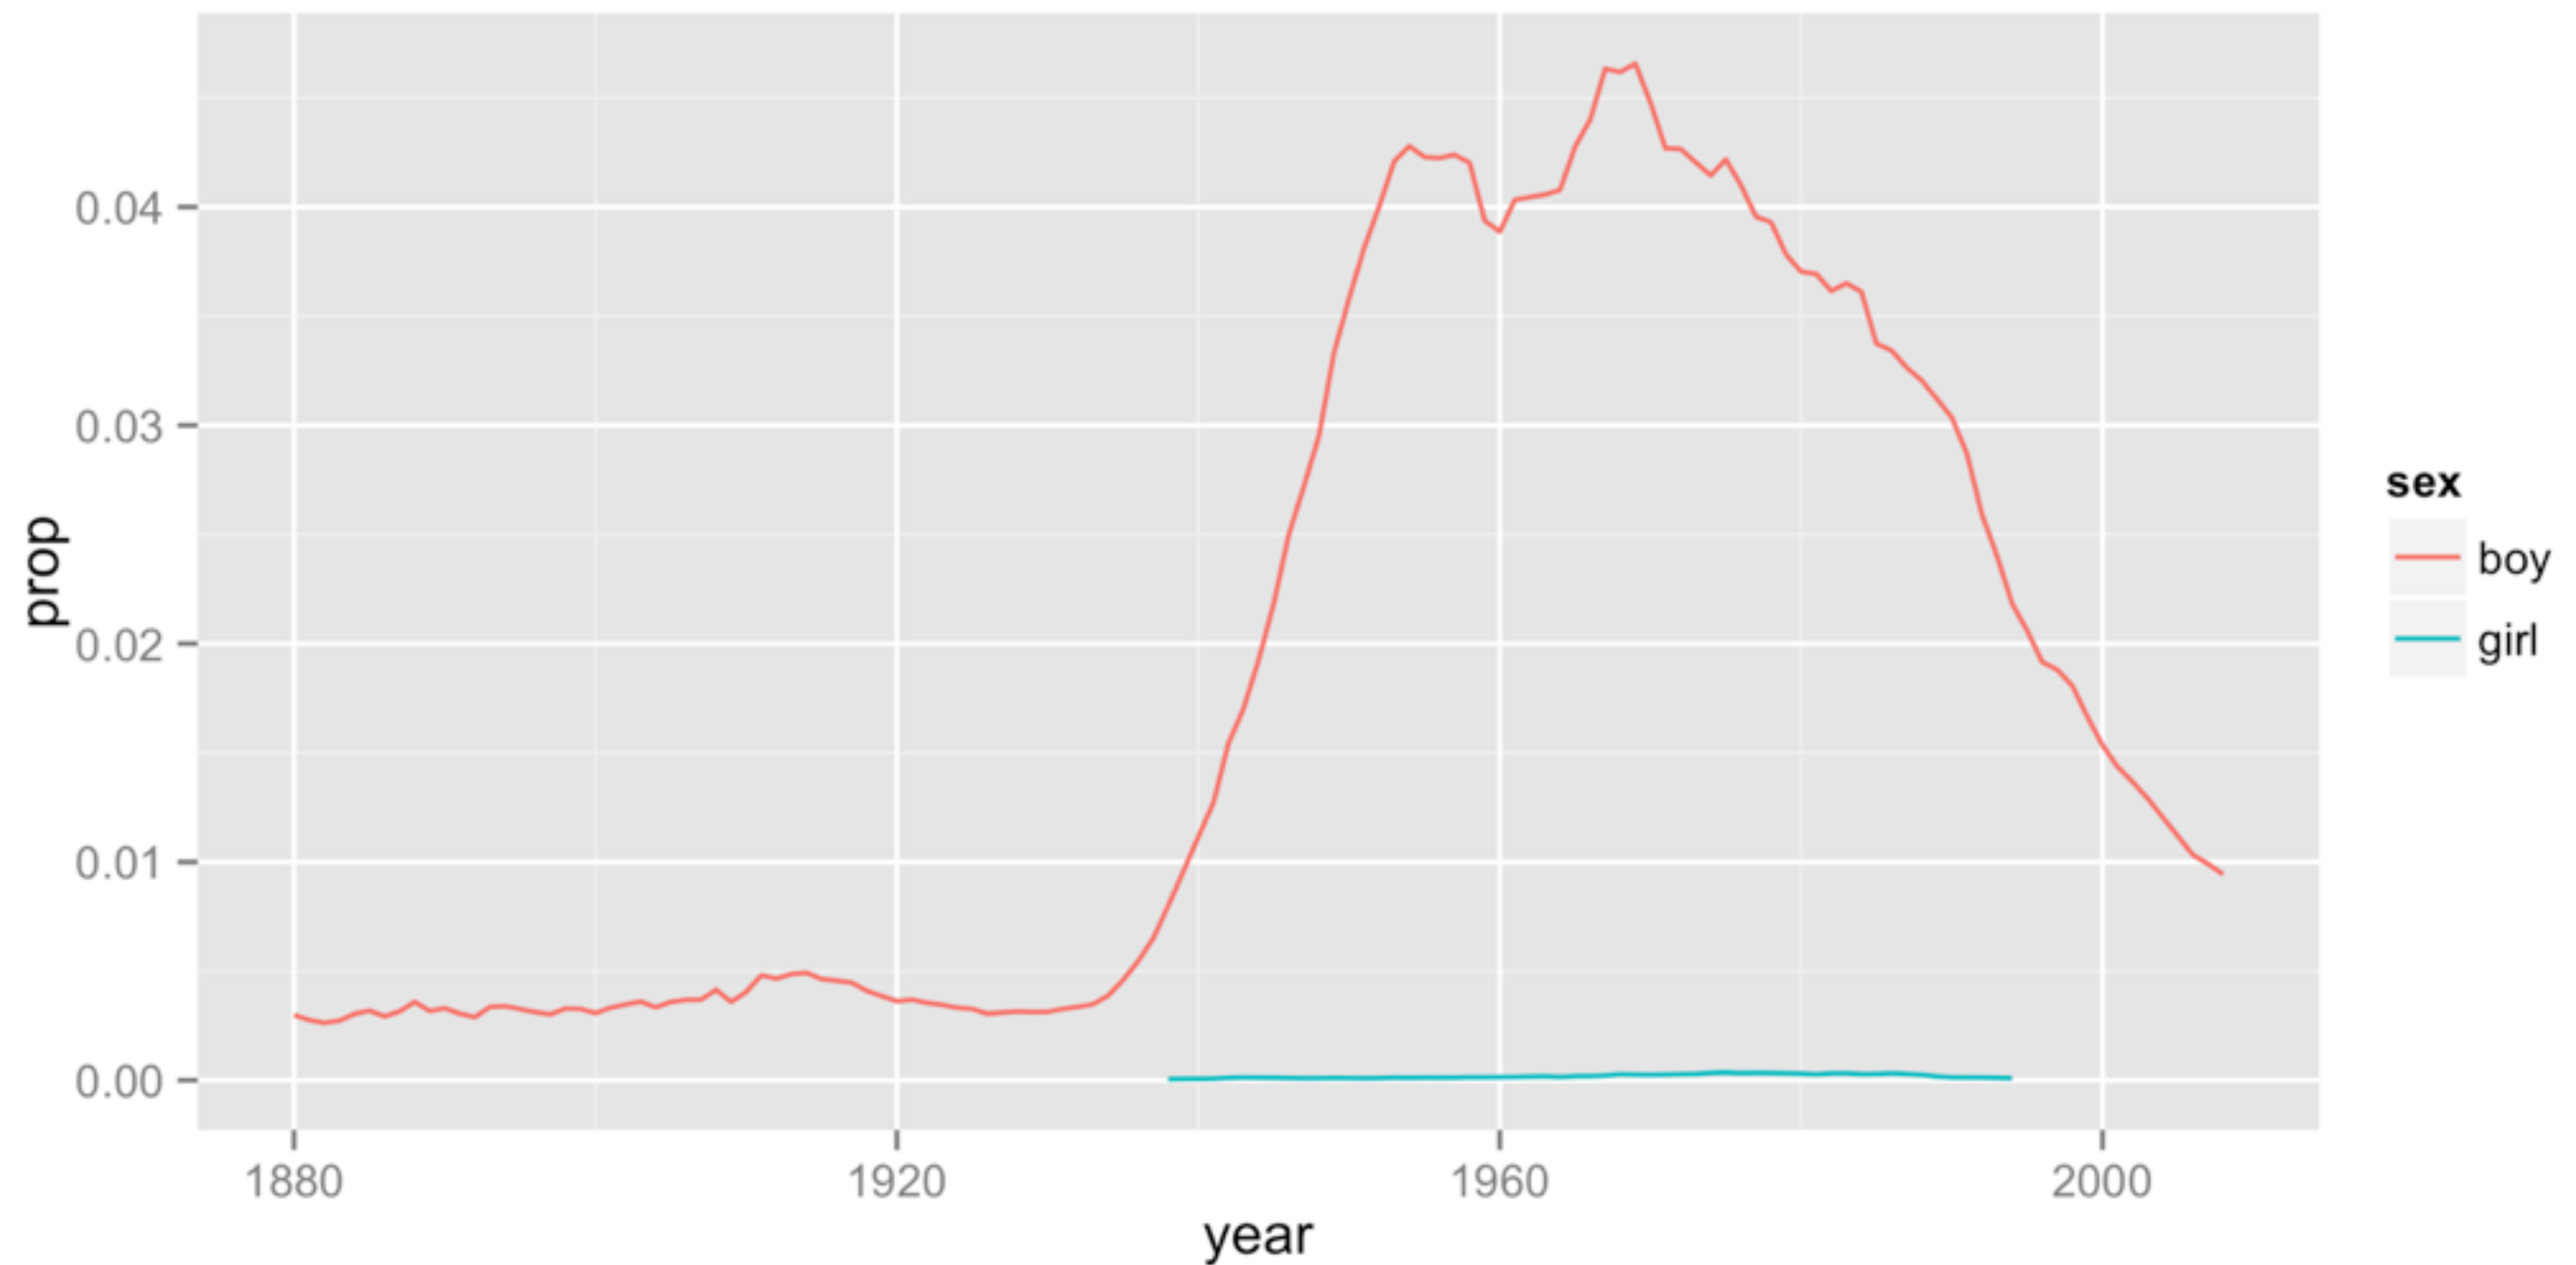

What is happening?



```
michael <- bnames[bnames$name == "Michael", ]  
qplot(year, prop, data = michael, geom = "line")
```

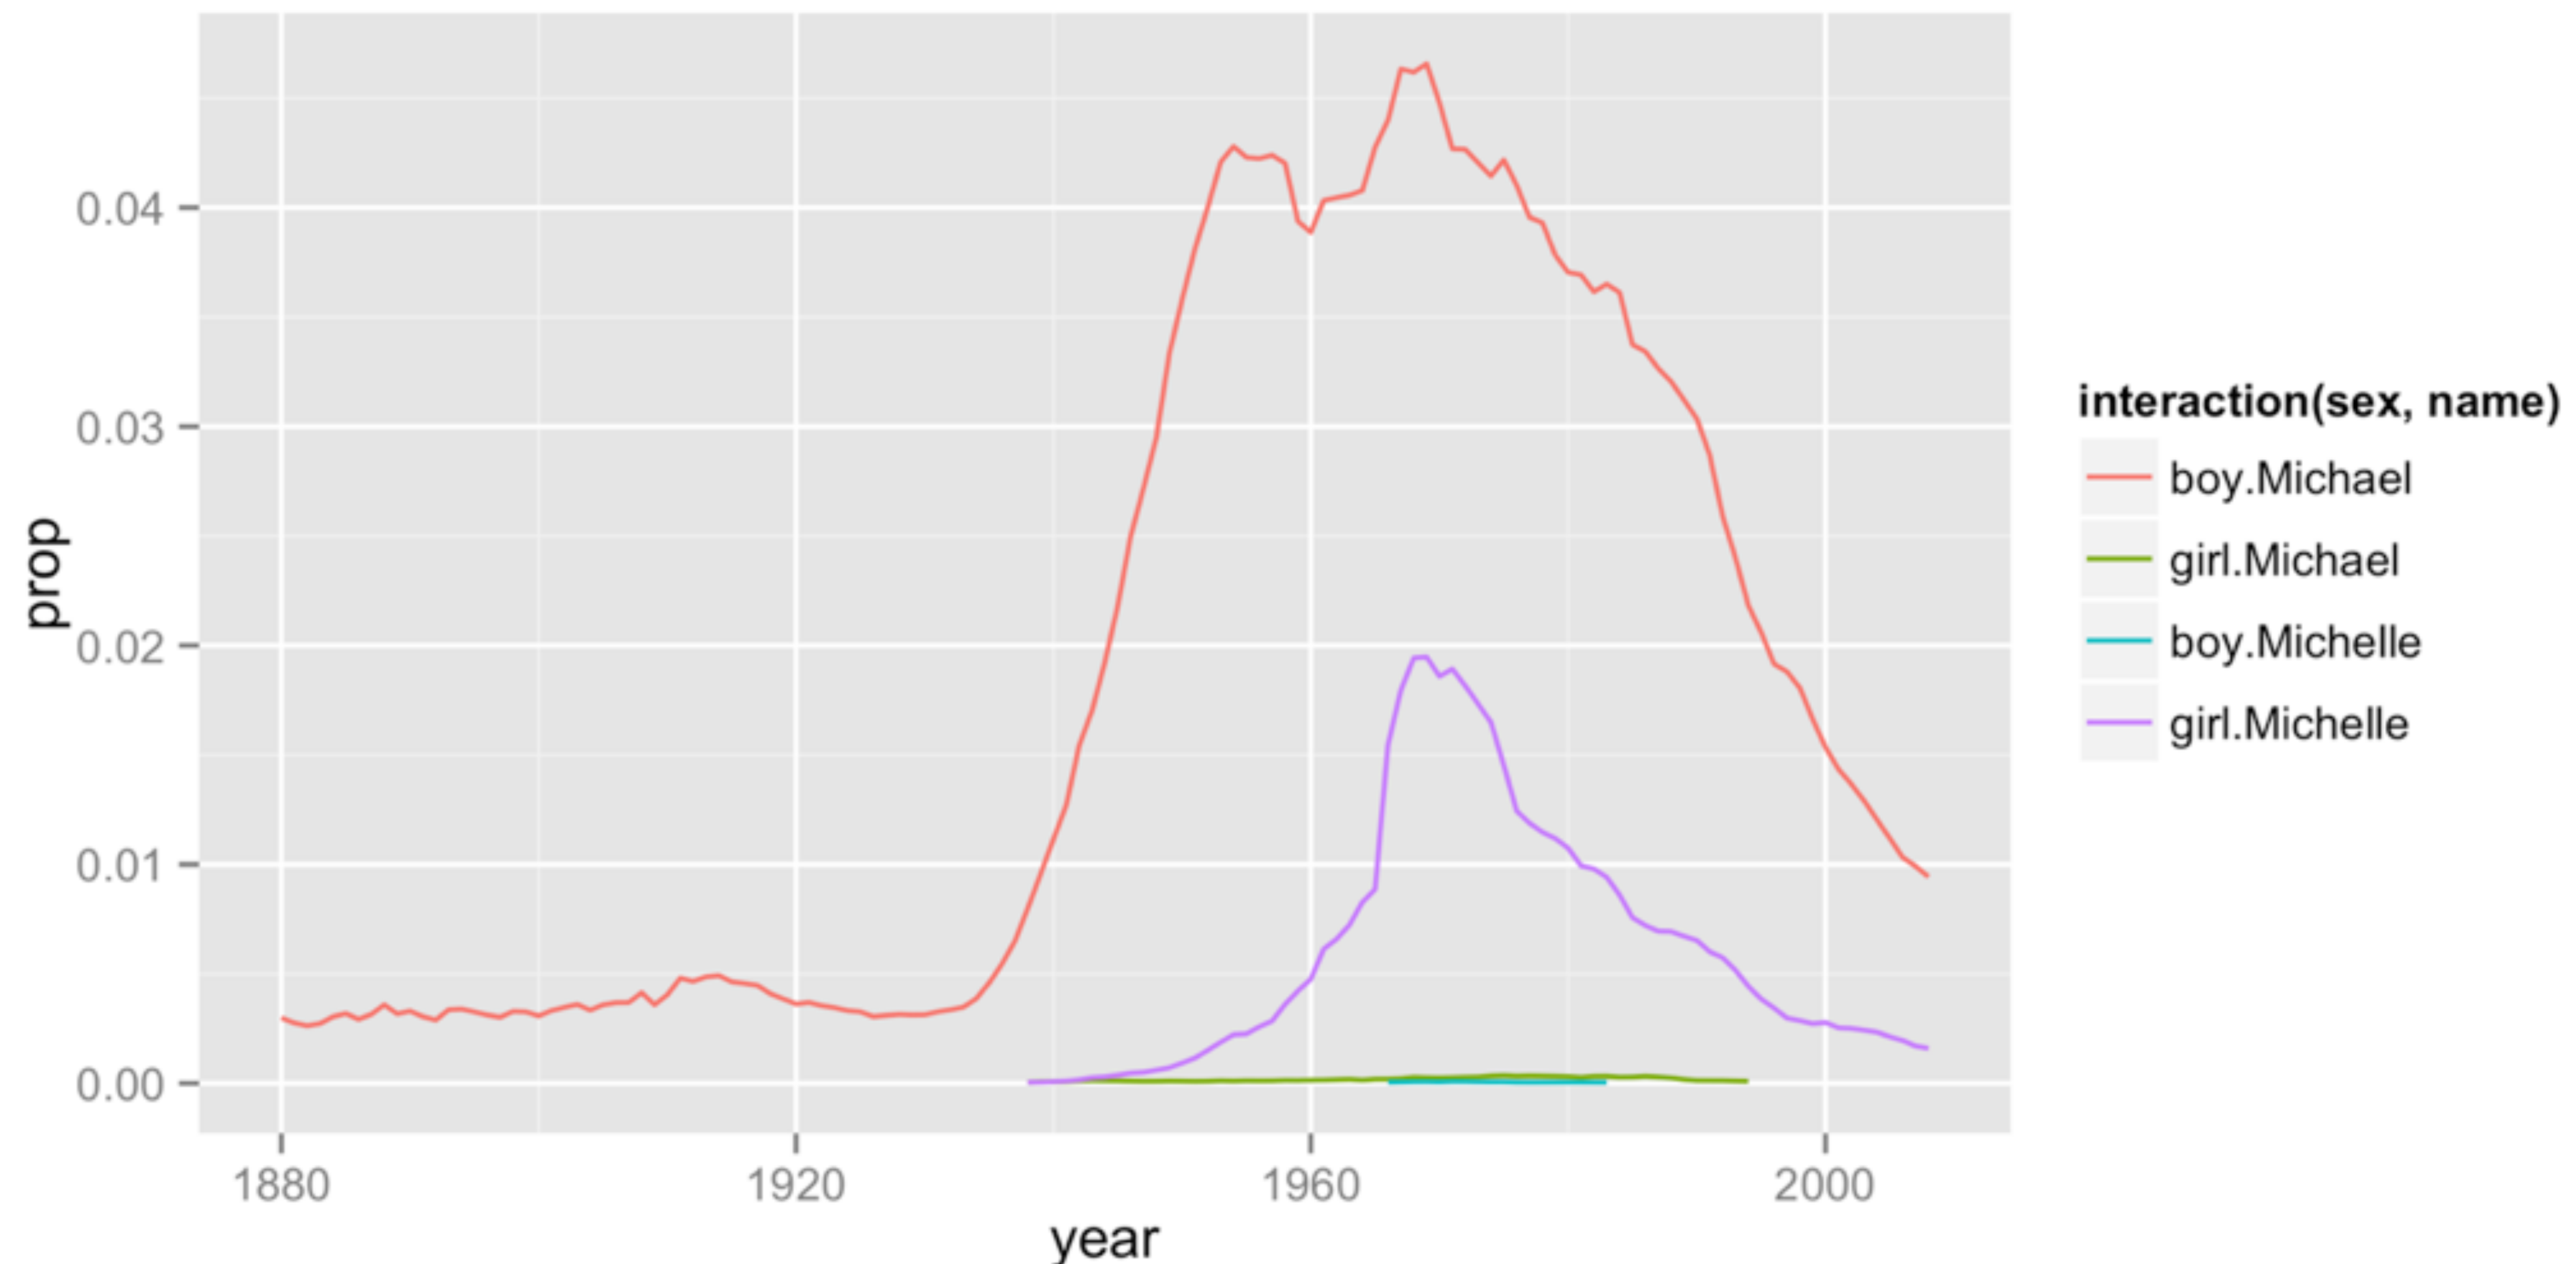


```
qplot(year, prop, data = michael, geom = "point")
```

```
qplot(year, prop, data = michael, geom = "line",  
color = sex)
```

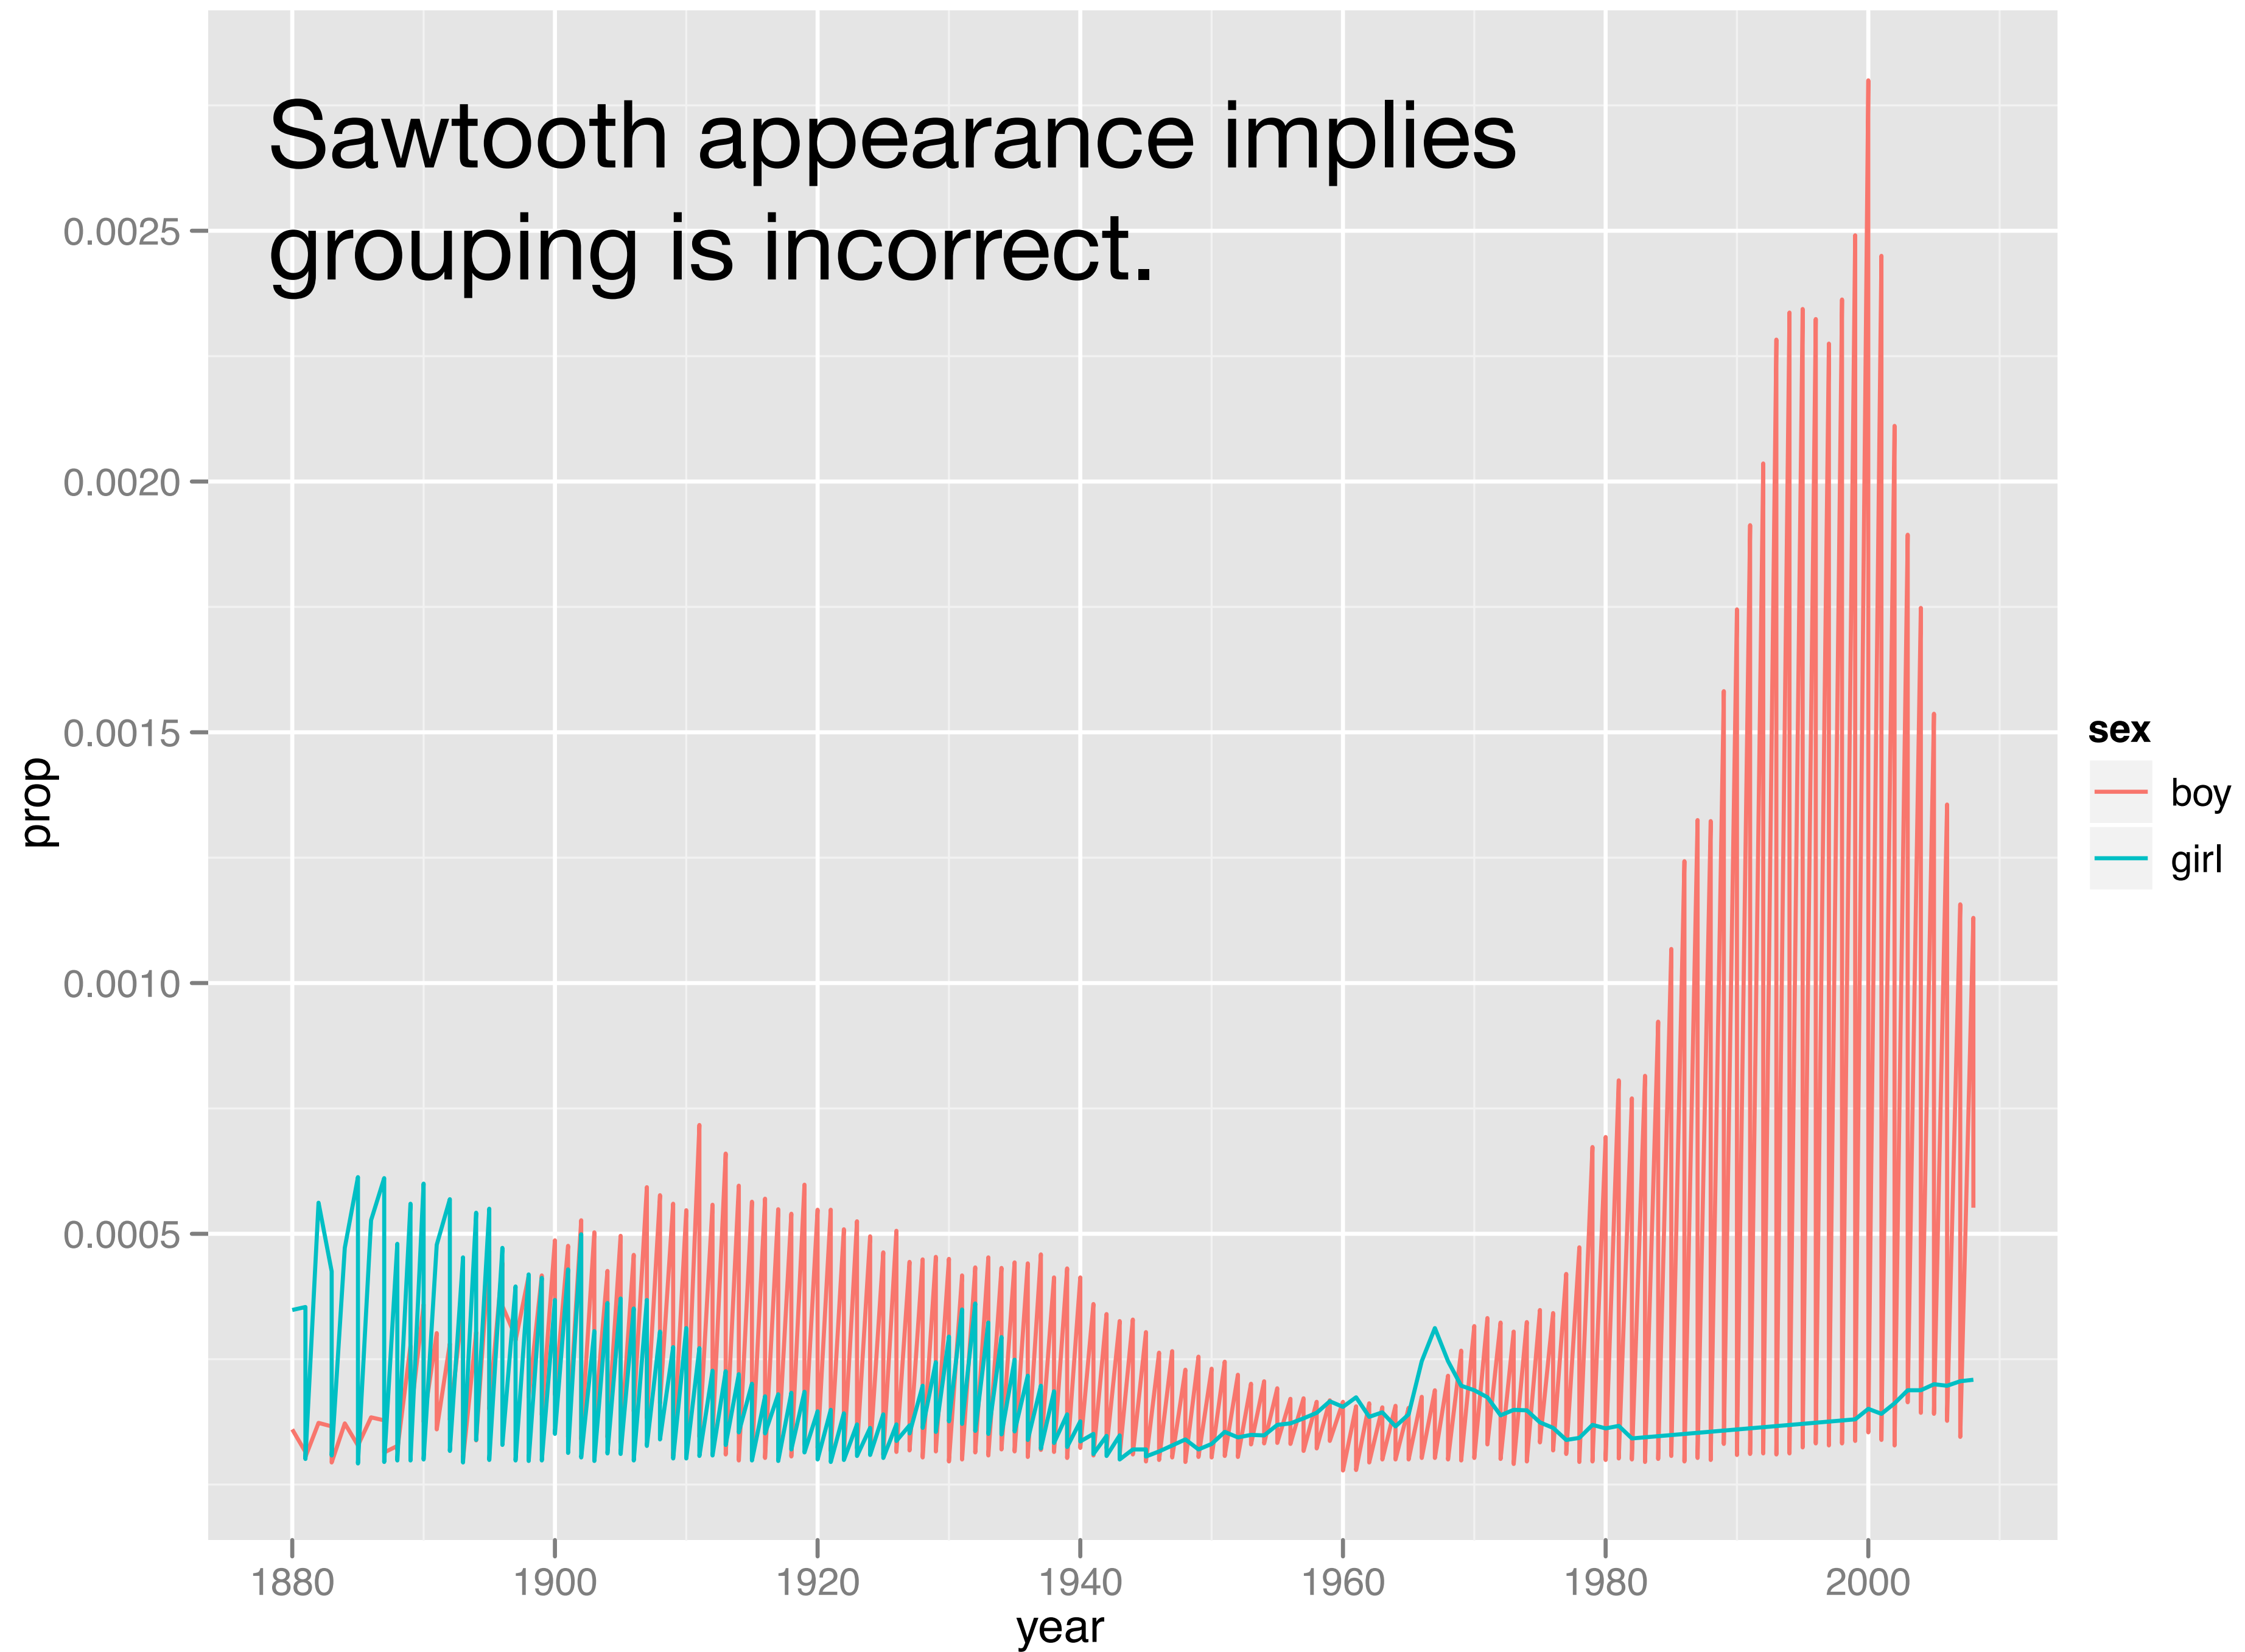
creates a different colored
line for each group of sex
(male, female)



```
michaels <- bnames[bnames$name == "Michael" |  
                  bnames$name == "Michelle", ]  
qplot(year, prop, data = michaels, geom = "line",  
       color = interaction(sex, name))
```

use interaction to group
on the combination of two
variables

Sawtooth appearance implies
grouping is incorrect.



dplyr

dplyr

An R package to manipulate data

- easier
- faster!

```
# install.packages("dplyr")  
library(dplyr)
```

filter

select

mutate

summarise

arrange

Structure

- First argument is a data frame
- Subsequent arguments say what to do with data frame
- Always return a data frame
- (Never modify in place)


```
tbl <- data.frame(  
  color = c("blue", "black", "blue", "blue", "black"),  
  value = 1:5)
```

tbl

color	value
blue	1
black	2
blue	3
blue	4
black	5



color	value
blue	1
blue	3
blue	4

```
filter(tbl, color == "blue")
```

tbl

color	value
blue	1
black	2
blue	3
blue	4
black	5



color	value
blue	1
blue	4

```
filter(tbl, value %in% c(1, 4))
```


Quiz

Use filter to find all of the girls names from
a year that ends in `00

```
filter(bnames, sex == "girl" & (year == 1900 | year ==  
2000))
```

tbl

color	value
blue	1
black	2
blue	3
blue	4
black	5



color
blue
black
blue
blue
black

```
select(tbl, color)
```

tbl

color	value
blue	1
black	2
blue	3
blue	4
black	5



value
1
2
3
4
5

```
select(tbl, -color)
```


- **filter:** pick **rows** by criteria
- **select:** pick **columns** by name

tbl

color	value
4	1
1	2
5	3
3	4
2	5



color	value
1	2
2	5
3	4
4	1
5	3

```
arrange(tbl, color)
```

tbl

color	value
4	1
1	2
5	3
3	4
2	5



color	value
5	3
4	1
3	4
2	5
1	2

```
arrange(tbl, desc(color))
```

Your turn

Reorder the rows from highest to lowest prop.

Which name was the most popular in a single year?

In what year was your name the most popular (hint: use the data set with just your name)


```
arrange(bnames, desc(prop)) # John in 1880  
arrange(garrett, desc(prop)) # 2000
```

tbl

color	value
blue	1
black	2
blue	3
blue	4
black	5



color	value	double
blue	1	2
black	2	4
blue	3	6
blue	4	8
black	5	10

```
mutate(tbl, double = 2 * value)
```

tbl

color	value
blue	1
black	2
blue	3
blue	4
black	5



color	value	double	quadruple
blue	1	2	4
black	2	4	8
blue	3	6	12
blue	4	8	16
black	5	10	20

```
mutate(tbl, double = 2 * value,  
       quadruple = 2 * double)
```

tbl

color	value
blue	1
black	2
blue	3
blue	4
black	5



total
15

```
summarise(tbl, total = sum(value))
```

tbl

color	value
blue	1
black	2
blue	3
blue	4
black	5



total	avg
15	3

```
summarise(tbl, total = sum(value),  
          avg = mean(value))
```


- **mutate:** add new variables
- **summarise:** reduce variables to values
- **arrange:** reorder rows

Your turn

With the data frame containing your name:

1. Add a new column to the data that changes the prop to a percentage
2. Create a summary that displays the min, mean, and max prop for your name.

```
mutate(garrett, perc = prop * 100)
```

```
summarise(garrett,  
  min = min(prop),  
  mean = mean(prop),  
  max = max(prop))
```

Summary

- **filter:** pick rows matching criteria
- **select:** pick columns by name
- **arrange:** reorder rows
- **mutate:** add new variables
- **summarise:** reduce variables to values

Joining data sets

Your turn

Why might prop be a bad way to compare names across different years?

Births

The number of boys and girls born
each year from 1880 - 2009

260 records ($2 * 130$)

Three variables: year, sex, births


```
head(bnames)
```

	year	name	prop	sex	sindex
1	1880	John	0.081541	boy	J500
2	1880	William	0.080511	boy	W450
3	1880	James	0.050057	boy	J520
4	1880	Charles	0.045167	boy	C642
5	1880	George	0.043292	boy	G620
6	1880	Frank	0.027380	boy	F652

```
head(births)
```

	year	sex	births
1	1880	boy	118405
2	1881	boy	108290
3	1882	boy	122034
4	1883	boy	112487
5	1884	boy	122745
6	1885	boy	115948

How would you combine these data sets?
Describe a strategy.

Joining datasets

name	instrument
John	guitar
Paul	bass
George	guitar
Ringo	drums
Stuart	bass
Pete	drums

+

name	band
John	T
Paul	T
George	T
Ringo	T
Brian	F

=

?

```
x <- data.frame(  
  name = c("John", "Paul", "George", "Ringo", "Stuart", "Pete"),  
  instrument = c("guitar", "bass", "guitar", "drums", "bass",  
    "drums"))
```

```
y <- data.frame(  
  name = c("John", "Paul", "George", "Ringo", "Brian"),  
  band = c("TRUE", "TRUE", "TRUE", "TRUE", "FALSE"))
```

x

y

name	instrument
John	guitar
Paul	bass
George	guitar
Ringo	drums
Stuart	bass
Pete	drums

+

name	band
John	T
Paul	T
George	T
Ringo	T
Brian	F

=

name	instrument	band
John	guitar	T
Paul	bass	T
George	guitar	T
Ringo	drums	T
Stuart	bass	NA
Pete	drums	NA

```
left_join(x, y, by = "name")
```


x**y**

name	instrument
John	guitar
Paul	bass
George	guitar
Ringo	drums
Stuart	bass
Pete	drums

+

name	band
John	T
Paul	T
George	T
Ringo	T
Brian	F

=

name	instrument	band
John	guitar	T
Paul	bass	T
George	guitar	T
Ringo	drums	T

```
inner_join(x, y, by = "name")
```

x**y**

name	instrument
John	guitar
Paul	bass
George	guitar
Ringo	drums
Stuart	bass
Pete	drums

+

name	band
John	T
Paul	T
George	T
Ringo	T
Brian	F

=

name	instrument
John	guitar
Paul	bass
George	guitar
Ringo	drums

```
semi_join(x, y, by = "name")
```

x

name	instrument
John	guitar
Paul	bass
George	guitar
Ringo	drums
Stuart	bass
Pete	drums

y

name	band
John	T
Paul	T
George	T
Ringo	T
Brian	F

+

=

name	instrument
Stuart	bass
Pete	drums

```
anti_join(x, y, by = "name")
```

Type	Action
left_join	Include all of x, and matching rows of y
inner_join	Include rows of x that appear in y, and matching rows of y
semi_join	Include rows of x that appear in y
anti_join	Include rows of x that <i>do not</i> appear in y

Your turn

Combine `bnames` with `births`, and then create a new column that shows the total number of babies born each year for each name.

```
bnames2 <- left_join(bnames, births, by = c("year",  
"sex"))
```

```
bnames2 <- mutate(bnames2, n = prop * births)  
bnames2
```

```
bnames2 <- mutate(bnames2, n = round(prop * births))  
bnames2
```


bnames2

```
##      year      name      prop sex soundex births      n
## 1  1880      John 0.081541 boy   J500  118405  9655
## 2  1880 William 0.080511 boy   W450  118405  9533
## 3  1880      James 0.050057 boy   J520  118405  5927
## 4  1880 Charles 0.045167 boy   C642  118405  5348
## 5  1880   George 0.043292 boy   G620  118405  5126
## 6  1880      Frank 0.027380 boy   F652  118405  3242
## 7  1880   Joseph 0.022229 boy   J210  118405  2632
## 8  1880   Thomas 0.021401 boy   T520  118405  2534
## 9  1880      Henry 0.020641 boy   H560  118405  2444
## 10 1880   Robert 0.020404 boy   R163  118405  2416
## .. ...      ...      ...      ...      ...      ...
```

Group wise operations

Total number of people per name

	name	total
1	Aaden	959
2	Aaliyah	39665
3	Aarav	219
4	Aaron	509464
5	Ab	25
6	Abigail	2682
7	Abb	16
8	Abbey	14348
9	Abbie	16622
10	Abbigail	6800

Do we have
enough
information to
calculate this?

Your turn

Calculate the total for a single name (e.g, your name).

Then devise a strategy for calculating the total for *every* name.

	name	total
1	Aaden	959
2	Aaliyah	39665
3	Aarav	219
4	Aaron	509464
5	Ab	25
6	Abigail	2682
7	Abb	16
8	Abbey	14348
9	Abbie	16622
10	Abbigail	6800

```
garrett <- filter(bnames2, name == "Garrett")  
sum(garrett$n)
```

Or

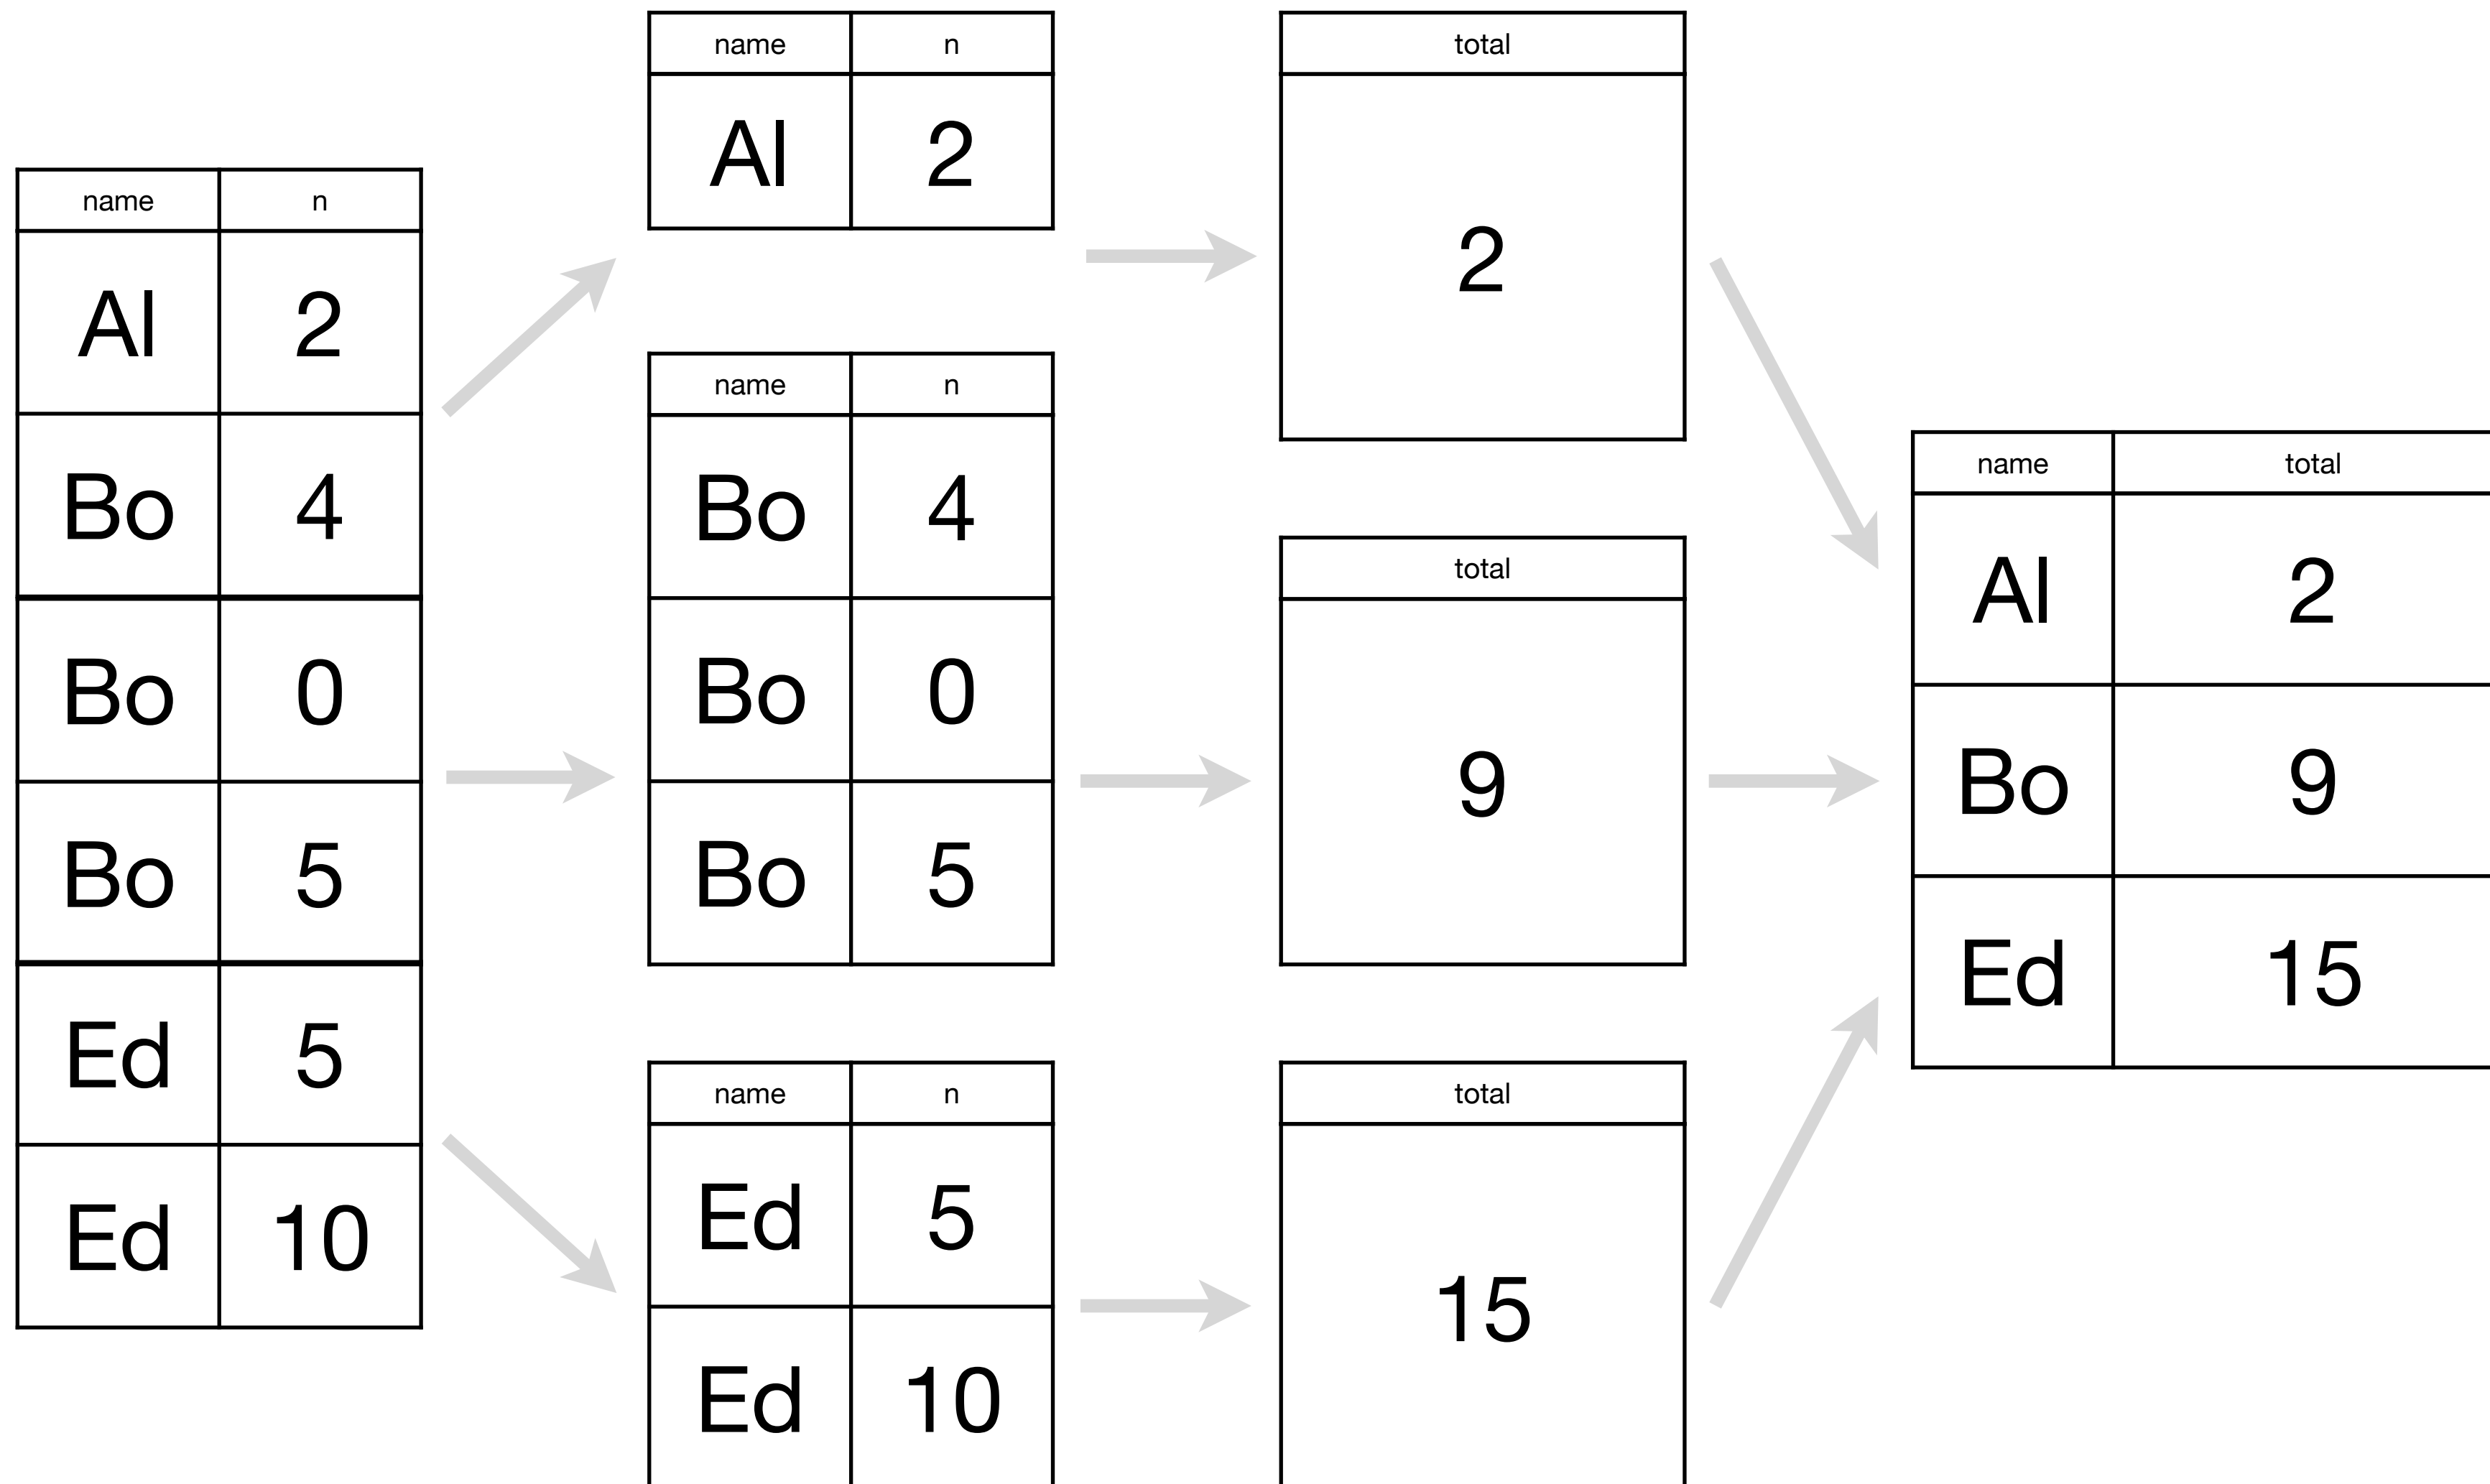
```
summarise(garrett, total = sum(n))
```

But how could we do this for every name?

Split

Apply

Combine



```
summarise(by_name, total = sum(n))
```

group_by


```
group_by(bnames2, name)
```

```
## Source: local data frame
```

```
## Groups: name
```

mutate, summarise, and
arrange will execute
groupwise on this variable

```
##   year   name   prop sex soundex births    n
## 1  1880   John 0.081541 boy   J500  118405 9655
## 2  1880 William 0.080511 boy   W450  118405 9533
## 3  1880   James 0.050057 boy   J520  118405 5927
## 4  1880 Charles 0.045167 boy   C642  118405 5348
## 5  1880  George 0.043292 boy   G620  118405 5126
## 6  1880   Frank 0.027380 boy   F652  118405 3242
## 7  1880 Joseph 0.022229 boy   J210  118405 2632
## 8  1880 Thomas 0.021401 boy   T520  118405 2534
## 9  1880   Henry 0.020641 boy   H560  118405 2444
## 10 1880 Robert 0.020404 boy   R163  118405 2416
## .. ... ..
```

tbl

color	value
blue	1
black	2
blue	3
blue	4
black	5



total
15

```
summarise(tbl, total = sum(value))
```

tbl

color	value
blue	1
black	2
blue	3
blue	4
black	5



color	total
blue	8
black	7

```
by_color <- group_by(tbl, color)
summarise(by_color, total = sum(value))
```

```
by_name <- group_by(bnames2, name)
totals <- summarise(by_name, total = sum(n))
totals
```

```
## Source: local data frame [6,782 x 2]
```

```
##      name  total
## 1   Aaden    959
## 2 Aaliyah 39665
## 3   Aarav   219
## 4   Aaron 509464
## 5     Ab    25
## 6 Abigail  2682
## 7    Abb    16
## 8   Abbey 14348
## ..      ...
```

tbl

tbl_df

tbl is a special case of data frame that can be manipulated more easily

```
bnames <- tbl_df(bnames)  
births <- tbl_df(births)
```

```
class(bnames)
```

```
## [1] "tbl_df"
```

```
"tbl"
```

```
"d"
```

Always start by making
your data tbls

R will always show just the part of the tbl that fits the console

```
Console ~/Dropbox (RStudio)/RStudio/rstudio-training/in-person-intro/Two-
> tbl_df(diamonds)
Source: local data frame [53,940 x 10]

  carat    cut color clarity depth table price
1  0.23   Ideal    E    SI2   61.5    55   326
2  0.21 Premium    E    SI1   59.8    61   326
3  0.23    Good    E    VS1   56.9    65   327
4  0.29 Premium    I    VS2   62.4    58   334
5  0.31    Good    J    SI2   63.3    58   335
6  0.24 Very Good  J    VVS2   62.8    57   336
7  0.24 Very Good  I    VVS1   62.3    57   336
8  0.26 Very Good  H    SI1   61.9    55   337
9  0.22    Fair    E    VS2   65.1    61   337
10 0.23 Very Good  H    VS1   59.4    61   338
.. ...      ...    ...    ...    ...    ...    ...
Variables not shown: x (dbl), y (dbl), z (dbl)
>
```

```
Console ~/Dropbox (RStudio)/RStudio/rstudio-training/
> tbl_df(diamonds)
Source: local data frame [53,940 x 10]

  carat    cut color clarity
1  0.23   Ideal    E    SI2
2  0.21 Premium    E    SI1
3  0.23    Good    E    VS1
4  0.29 Premium    I    VS2
5  0.31    Good    J    SI2
6  0.24 Very Good  J    VVS2
7  0.24 Very Good  I    VVS1
8  0.26 Very Good  H    SI1
9  0.22    Fair    E    VS2
10 0.23 Very Good  H    VS1
.. ...      ...    ...    ...
Variables not shown: depth (dbl),
  table (dbl), price (int), x
  (dbl), y (dbl), z (dbl)
> |
```

Use View() to see more

Interactions

Interactions

Use multiple variables to create groups based on the interaction of variables

```
group_by(bnames2, name, sex)
```

```
## Source: local data frame
```

```
## Groups: name, sex
```

dplyr will treat each unique combination of these values as a separate group

```
##   year   name   prop sex soundex births  n
## 1  1880   John 0.081541 boy   J500  118405 9655
## 2  1880 William 0.080511 boy   W450  118405 9533
## 3  1880   James 0.050057 boy   J520  118405 5927
## 4  1880 Charles 0.045167 boy   C642  118405 5348
## .. ... ..
```

Interactions

```
by_name <- group_by(bnames2, name)
group_by(by_name, sex)
## Source: local data frame [258,000 x 7]
## Groups: name, sex
```

```
##   year      name      prop sex soundex births    n
## 1  1880     John 0.081541 boy   J500  118405 9655
## 2  1880 William 0.080511 boy   W450  118405 9533
## 3  1880    James 0.050057 boy   J520  118405 5927
## 4  1880  Charles 0.045167 boy   C642  118405 5348
## .. ... ..
```

```
name_sex <- group_by(bnames2, name, sex)
totals2 <- summarise(name_sex, total = sum(n))
totals2
```

```
## Source: local data frame [7,455 x 3]
```

```
## Groups: name
```

```
##      name  sex  total
## 1   Aaden  boy   959
## 2 Aaliyah girl 39665
## 3   Aarav  boy   219
## 4   Aaron  boy 508094
## 5   Aaron girl  1370
## 6     Ab   boy    25
## 7 Abagail girl  2682
## ..      ...      ...
```

summarise returns a data frame that has one less level of grouping

“Unwrap” groups with summarise


```
name_sex <- group_by(bnames2, name, sex)
totals2 <- summarise(name_sex, total = sum(n))
```

```
totals2
```

```
## Source: local data frame [7,455 x 3]
```

```
## Groups: name
```

```
##      name  sex  total
## 1   Aaden  boy   959
## 2 Aaliyah girl 39665
## 3   Aarav  boy   219
## 4   Aaron  boy 508094
## 5   Aaron girl  1370
## 6     Ab   boy    25
## 7 Abigail girl  2682
## ..      ...  ...   ...
```



```
summarise(totals2, total = sum(total))
```

```
Source: local data frame [6,782 x 2]
```

```
      name  total
1   Aaden   959
2 Aaliyah 39665
3   Aarav   219
4   Aaron 509464
5     Ab    25
6 Abigail  2682
7     Abb   16
8   Abbey 14348
9   Abbie 16622
10 Abbigail 6800
..      ...   ...
```

Ungroup

Use `ungroup` to remove group specifications.

```
by_name_sex <- group_by(bnames2, name, sex)  
ungroup(by_name_sex)
```

```
## Source: local data frame [258,000 x 7]
```

```
##   year      name      prop sex soundex births      n  
## 1  1880     John 0.081541 boy   J500  118405  9655  
## 2  1880 William 0.080511 boy   W450  118405  9533  
## 3  1880    James 0.050057 boy   J520  118405  5927  
## 4  1880 Charles 0.045167 boy   C642  118405  5348  
## .. ... ..
```

Summary functions

- `min(x)`, `median(x)`, `max(x)`,
`quantile(x, p)`
- `n()`, `nth()`, `n_distinct()`
- `sd(x)`, `var(x)`, `iqr(x)`, `mad(x)`

Practice

Your turn

Calculate the total number of babies in each soundex

```
by_soundex <- group_by(bnames2, soundex)  
stotals <- summarise(by_soundex, total = sum(n))  
stotals
```

```
## Source: local data frame [1,392 x 2]
```

```
##      soundex  total  
## 1      A000      11  
## 2      A100 193837  
## 3      A120  15652  
## 4      A124 256458  
## 5      A130      11  
## 6      A134   5181  
## 7      A135    901  
## 8      A140  30815  
## ..      ...      ...
```

```
# What is the most popular sound?  
arrange(stotals, desc(total))
```

What is the most popular sound?

arrange(stotals, desc(total))

Source: local data frame [1,392 x 2]

##	soundex	total
## 1	J500	9991737
## 2	M240	5823791
## 3	M600	5553703
## 4	J520	5524958
## 5	R163	5047182
## 6	W450	4116109
## 7	C623	4016919
## 8	J200	3859240
## 9	D500	3774287
## 10	D130	3506995
##

What is the most popular sound?

```
j500 <- filter(bnames, soundex == "J500")
unique(j500$name)
```

```
## [1] "John"      "Jim"       "Juan"      "Jimmie"    "Johnnie"   "Johnny"
## [7] "Johnnie"   "Jean"      "June"      "Jonah"     "Jennie"    "Jimmy"
## [13] "Johnny"    "Jonnie"    "Johney"    "Jamie"     "Jon"       "Joan"
## [19] "Jan"       "Jame"      "Jaime"     "Jamey"     "Jaimie"    "Jammie"
## [25] "Jayme"     "Juwan"     "Johan"     "Jaheim"    "Jahiem"    "Jaheem"
## [31] "Jane"      "Janie"     "Johanna"   "Joanna"    "Jannie"    "Jenny"
## [37] "Jeanne"    "Johannah" "Juana"     "Junie"     "Jinnie"    "Jeanie"
## [43] "Jeannie"   "Junia"     "Janey"     "Jeane"     "Joanne"    "Joann"
## [49] "Jayne"     "Jana"      "Janna"     "Jann"      "Joni"      "Joanie"
## [55] "Jeanna"    "Jami"      "Johnna"    "Jeana"     "Jonna"     "Jena"
## [61] "Jenni"     "Jenna"     "Janae"     "Jaimee"    "Janay"     "Joana"
## [67] "Janiya"    "Johana"    "Jamy"      "Janiyah"   "Janiah"    "Jamiya"
```


Your turn


Calculate the total number of boys and the total number of girls for each year

```
year_sex <- group_by(bnames2, year, sex)  
ytotals <- summarise(year_sex, births = sum(n))  
ytotals
```

Source: local data frame [258 x 3]

Groups: year

	year	sex	births
1	1880	boy	110207
2	1880	girl	91227
3	1881	boy	100763
4	1881	girl	92204
5	1882	boy	113194
6	1882	girl	107713
7	1883	boy	104487
..



```
qplot(year, total, data = ytotals, geom = "line", color = sex)
```

Your turn

Calculate the rank of each name within each year (and within each gender).


Hint: create a new variable equal to `rank(desc(prop))`

```
year_sex <- group_by(bnames2, year, sex)  
ranks <- mutate(year_sex, rank = rank(desc(prop)))  
ranks
```

```
## Source: local data frame [258,000 x 8]
```

```
## Groups: year, sex
```

##	year	name	prop	sex	soundex	births	n	rank
## 1	1880	John	0.081541	boy	J500	118405	9655	1
## 2	1880	William	0.080511	boy	W450	118405	9533	2
## 3	1880	James	0.050057	boy	J520	118405	5927	3
## 4	1880	Charles	0.045167	boy	C642	118405	5348	4
## 5	1880	George	0.043292	boy	G620	118405	5126	5
## 6	1880	Frank	0.027380	boy	F652	118405	3242	6
## 7	1880	Joseph	0.022229	boy	J210	118405	2632	7
##



```
garrett <- filter(ranks, name == "Garrett")  
qplot(year, -rank, data = garrett, geom = "line")
```

What names were ranked number one?

```
ones <- filter(ranks, rank == 1)
```

```
ones <- select(ones, year, name, sex)
```

```
ones
```

```
## Source: local data frame [258 x 3]
```

```
## Groups: year, sex
```

```
##      year name sex
```

```
## 1  1880 John boy
```

```
## 2  1881 John boy
```

```
## 3  1882 John boy
```

```
## 4  1883 John boy
```

```
## 5  1884 John boy
```

```
## 6  1885 John boy
```

```
## 7  1886 John boy
```

```
## 8  1887 John boy
```

```
## 9  1888 John boy
```

```
## 10 1889 John boy
```

```
..    ..    ..    ..
```


What names were ranked number one?

library(reshape2)

dcast(ones, year ~ sex, value = "name")

##	year	boy	girl
## 1	1880	John	Mary
## 2	1881	John	Mary
## 3	1882	John	Mary
## 4	1883	John	Mary
## 5	1884	John	Mary
## 6	1885	John	Mary
## 7	1886	John	Mary
## 8	1887	John	Mary
## 9	1888	John	Mary
## 10	1889	John	Mary

**Where
next?**

Other data sources

dplyr creates a consistent syntax for manipulating data from various sources

- data frame, data table, data cube
- PostgreSQL, Greenplum, redshift
- MySQL, MariaDB
- SQLite
- MonetDB, BigQuery
- *Oracle, SQL Server*

Read dplyr's built in vignettes to learn more

```
browseVignettes(package = "dplyr")
```